

MPI Distributed Systems Assignment Report

Question 1: Program for Finding K Nearest Neighbours

The program follows these steps:

1. Initialization:

- The program initializes MPI, determines the rank of each process, and the total number of processes.

2. Input Handling:

- The root process (rank 0) reads input values: the number of points in sets P (N) and Q (M), and the value of K .
- Points in sets P and Q are also read by the root process.
- The values of N , M , and K are broadcast to all other processes using `MPI_Bcast`.

3. Data Distribution:

- The points in set Q are divided among the processes. Each process receives a chunk of points (approximately M/size points, with some processes possibly getting one extra point if M is not divisible by the number of processes).
- The root process sends the corresponding subset of Q to each process using `MPI_Scatterv`.

4. Local Computation:

- Each process computes the K closest points in P for each point in its assigned subset of Q .
- For each point in the local subset of Q , the program calculates the Euclidean distance to all points in P .
- A max-heap is used to maintain the closest K points efficiently.
- The result for each point in Q is stored locally.

5. Gathering Results:

- Each process sends its local results back to the root process using `MPI_Gatherv`.
- The root process collects all results and prints them.

6. Output:

- The root process prints the K closest points for each point in Q in order.

Time Complexity

1. Broadcasting P and Q :

- Broadcasting P : $O(N)$ per process for receiving.
- Broadcasting Q for scatter: $O(M/size)$ per process for receiving.

2. Local Computation:

- Each process handles $M/size$ points from Q . For each point in Q , calculating the distance to all N points in P takes $O(N)$.
- Maintaining a max-heap of size K for the closest points takes $O(\log K)$ per insertion.
- Total time per process: $O((M/size) * N * \log K)$.

3. Gathering Results:

- Gathering results from all processes has a cost of $O(M)$ in total for the root process.
- **Total Time Complexity:** $O(M * N * \log K / size)$.

Space Complexity

1. Memory for Point Storage:

- The entire P set of N points is stored on all processes: $O(N)$.
- Each process stores approximately $M/size$ points from Q : $O(M/size)$.

2. Local Results:

- Each process stores results for its local subset of Q , consuming $O(K * M/size)$ space.

3. Global Results:

- The root process stores all results, requiring $O(K * M)$ space.
- **Total Space Complexity:**
 - Per process: $O(N + K * M/size)$.
 - On the root process (rank 0): $O(N + K * M)$.

Question 2: MPI Program for Julia Set Calculation

The program performs the following steps:

1. Initialization:

- MPI is initialized, and each process retrieves its rank and the total number of processes ($size$).

2. Input Handling:

- The root process (rank 0) reads input values, including the dimensions of the grid (N, M), the maximum number of iterations (K), and the constant complex number $c = c_real + c_img*i$ used in the Julia set formula.

- These values are broadcasted to all processes using `MPI_Bcast`.

3. Grid Partitioning:

- The grid is partitioned row-wise among the processes. Each process is responsible for computing the Julia set values for a subset of rows. The number of rows each process handles is determined by `rows_per_proc`, with a remainder handled by some processes to ensure all rows are covered.

4. Local Computation:

- Each process computes the Julia set for its assigned rows.
- The grid points are mapped to complex numbers, $z_0 = x + yi$, within the range $[-1.5, 1.5]$ in both real and imaginary parts.
- For each grid point, the Julia set membership is determined by iterating the complex function $z = z * z + c$ up to K times. If the magnitude of z exceeds a threshold T , the point is not part of the Julia set.

5. Gathering Results:

- Each process sends its computed results to the root process using `MPI_Gatherv`. This function gathers the computed data from all processes and assembles them into the final result.
- The root process collects the results into a global grid that represents the entire Julia set.

6. Output:

- The root process prints the results, where each grid point is represented by either 1 (indicating the point is within the Julia set) or 0 (indicating the point is outside the set).

Time Complexity

1. Broadcasting Input:

- Broadcasting N , M , K , and the complex number c takes $O(1)$ time for the root process, but all processes need to receive this data, making it $O(1)$ per process.

2. Local Computation:

- Each process computes the Julia set for approximately N/size rows. For each point in the grid, computing the Julia set takes up to K iterations, resulting in:
- **Time Complexity per process:** $O((N/\text{size}) * M * K)$.

3. Gathering Results:

- Gathering the computed results at the root process has a cost of $O(N * M)$ for the root process as it must collect data from all processes.
- **Total Time Complexity:** The overall time complexity is dominated by the local computation, leading to:

$$\blacksquare O(N * M * K / \text{size}).$$

Space Complexity

1. Memory for Grid Storage:

- Each process stores approximately N/size rows of the grid, consuming $O((N/\text{size}) * M)$ space.

2. Global Result Storage:

- The root process stores the entire grid, which requires $O(N * M)$ space.

3. Communication Buffers:

- Each process maintains buffers for sending and receiving data, which add an additional $O(M)$ space for the row data.
 - Total Space Complexity:**
 - Per process: $O((N/\text{size}) * M)$.
 - On the root process (rank 0): $O(N * M)$ for the final output grid.
-

Question 3: Program for Finding the Inverse of a Matrix

The program follows these steps:

1. Initialization:

- The program initializes MPI, determines the rank of each process, and the total number of processes.

2. Input Handling:

- The root process (rank 0) reads the size of the matrix (N) and the matrix (A) itself.
- The value of (N) is broadcast to all other processes using `MPI_Bcast`.

3. Data Distribution:

- The matrix (A) is divided into rows, and each process receives a portion of these rows. The number of rows assigned to each process is approximately (N/size) , with some processes possibly receiving an additional row if (N) is not divisible by the number of processes.
- The rows of the matrix (A) are scattered to each process using `MPI_Scatterv`. Each process also augments its portion of the matrix with the corresponding identity matrix rows.

4. Local Computation:

- Each process performs Gaussian elimination on its local rows, normalizing the pivot row and broadcasting it to all other processes.
- Each process uses the broadcasted pivot row to eliminate the corresponding elements in its own rows, ensuring the matrix is reduced to row-echelon form.

5. Gathering Results:

- Once all rows are reduced, each process sends its portion of the inverse matrix back to the root process using `MPI_Gatherv`.
- The root process collects all the portions to reconstruct the complete inverse matrix.

6. Output:

- The root process prints the inverse matrix with exactly two decimal places for each element.

Time Complexity

1. Broadcasting (N) and Scattering (A):

- Broadcasting (N): $O(1)$ per process.
- Scattering (A): $O(N^2/\text{size})$ per process for receiving.

2. Gaussian Elimination:

- Each process handles approximately (N/size) rows. The time to normalize the pivot row and eliminate elements across all rows is dominated by the row-wise operations.
- **Overall time complexity for Gaussian elimination:** $O(N^3/\text{size})$.

3. Gathering Results:

- Gathering the inverse matrix from all processes: $O(N^2)$ for the root process.
- **Total Time Complexity:** $O(N^3/\text{size})$.

Space Complexity

1. Memory for Matrix Storage:

- Each process stores its portion of (A), which is approximately $O(N^2/\text{size})$.
- Each process also needs to store the identity matrix portion for augmentation: $O(N^2/\text{size})$.

2. Local Results:

- The inverse matrix's portion: $O(N^2/\text{size})$ per process.

3. Global Results:

- The root process stores the entire inverse matrix: $O(N^2)$.
 - **Total Space Complexity:**
 - Per process: $O(N^2/\text{size})$.
 - On the root process (rank 0): $O(N^2)$.
-

Question 4: Program for Parallel Matrix Chain Multiplication Problem

The program follows these steps:

1. Initialization:

- The program initializes MPI, determines the rank of each process, and the total number of processes.

2. Input Handling:

- The root process (rank 0) reads the number of matrices (N) and their dimensions from the input.
- The dimensions array is broadcast to all other processes using `MPI_Bcast`.

3. Data Distribution:

- The dynamic programming table for matrix chain multiplication is computed in parallel. The subproblems are distributed across the processes, with each process handling a subset of the table.
- Each process computes the cost of multiplying a specific range of matrices, based on its rank.

4. Local Computation:

- Each process calculates the minimum number of scalar multiplications required for its assigned portion of the matrix chain.
- The results from each step are synchronized using `MPI_Barrier` and broadcast to all processes to ensure consistency.

5. Gathering Results:

- The final minimum multiplication cost is gathered from all processes and collected by the root process.
- The root process then prints the result.

Time Complexity

1. Broadcasting Dimensions:

- Broadcasting the dimensions array: $O(N)$ per process.

2. Dynamic Programming Table Calculation:

- Each process handles a portion of the dynamic programming table, with the computation taking $O(N^3 / \text{size})$ per process.

3. Synchronization and Broadcasting Results:

- Synchronization overhead due to barriers: $O(N^2 \log \text{size})$ overall.
- **Total Time Complexity:** $O(N^3 / \text{size})$.

Space Complexity

1. Memory for Storing Dimensions:

- Each process stores the dimensions array: $O(N)$.

2. Dynamic Programming Table:

- Each process maintains a portion of the table: $O(N^2/\text{size})$.

3. Global Results:

- The root process stores the entire table for final output: $O(N^2)$.

- **Total Space Complexity:**

- Per process: $O(N^2/\text{size})$.
 - On the root process (rank 0): $O(N^2)$.
-

Question 5: MPI Program for Parallel Prefix Sum Calculation

The program performs the following steps:

1. Initialization:

- MPI is initialized, and each process retrieves its rank (`rank`) and the total number of processes (`size`).

2. Input Handling:

- The root process (rank 0) reads the input data, including the size of the sequence `n` and the sequence of `n` double values.
- These values are broadcasted to all processes using `MPI_Bcast`.

3. Grid Partitioning:

- The sequence is divided into `size` parts, one for each process.
- The division accounts for the possibility that `n` may not be perfectly divisible by `size`, so some processes may handle one additional element (`remainder`).

4. Local Computation:

- Each process computes the prefix sum for its assigned portion of the sequence.
- For each element in its local array, the prefix sum is calculated iteratively by adding the current element to the sum of all previous elements in the local array.

5. Gathering Local Sums:

- Each process calculates the last element of its local prefix sum and sends it to the root process using `MPI_Gather`.

- The root process uses these last elements to compute offsets for each process, which are then broadcasted back to all processes using `MPI_Bcast`.

6. Adjusting Local Prefix Sums:

- Each process adjusts its local prefix sum by adding the offset it received. This ensures that the prefix sum is correctly accumulated across all processes.

7. Gathering Results:

- The adjusted local prefix sums from all processes are gathered at the root process using `MPI_Gatherv`.
- The root process assembles the final prefix sum sequence and prints it.

Time Complexity

1. Broadcasting Input:

- Broadcasting the input size and data takes $O(1)$ time for the root process and $O(1)$ for each receiving process.

2. Local Prefix Sum Calculation:

- Each process computes the prefix sum for its portion of the sequence.
- **Time Complexity per process:** $O(n/size)$.

3. Gathering Last Elements:

- The root process gathers the last elements from all processes, which takes $O(size)$ time for the root process.

4. Broadcasting Offsets:

- The offsets are broadcasted to all processes, which takes $O(size)$ time.

5. Adjusting and Gathering Results:

- Adjusting local prefix sums and gathering the final results take $O(n)$ time for the root process.
- **Total Time Complexity:** The total time complexity is $O(n/size)$, dominated by the local prefix sum calculation.

Space Complexity

1. Memory for Local Arrays:

- Each process stores its portion of the sequence, which requires $O(n/size)$ space.

2. Global Result Storage:

- The root process stores the entire prefix sum sequence, requiring $O(n)$ space.

3. Communication Buffers:

- Additional space is used for communication buffers, such as `sendcounts`, `displacements`, and `offsets`, which add $O(\text{size})$ space.
- **Total Space Complexity:**
 - Per process: $O(n/\text{size})$.
 - On the root process (rank 0): $O(n)$ for the final output sequence.